

MEASURING THE IMPACT OF AI EDUCATION

Ezra Savard • Netflix • Enterprise AI Summit

April 2026



THE PROBLEM

By late 2025, Netflix had **60%+ WAU** on agentic coding tools

But the gap between heavy users and everyone else was **widening**

WORKSHOP DESIGN

4.5-hour small-group hands-on workshop

Bring your **real work**

Two things: brainstorm/plan with AI, and
take code to production

Minimal Claude Code setup — **teach
principles, not tool tricks**

23 sessions, 140+ engineers, 10 teachers by
end of Q1

STUDY DESIGN

Workshop group: **~90 engineers across 23 sessions** (Dec 2025 – Mar 2026)

Control group: **~3,500 engineers** — same tools, no workshop

Almost all enrolled by managers in team sessions, **not self-selected**

WHAT WE MEASURED

Daily token consumption

PR metrics — count, lines changed

Self-reported confidence surveys

Qualified Token Days (QTDs) — days with
≥4M tokens consumed

Line Cook = ≥3 QTDs/week = tool embedded
in daily workflow

KNOWN LIMITATIONS

Not an RCT — quasi-experiment with natural control

Selection bias: the org was AI-forward (we funded this), eager EMs booked earlier

Workshop group starts slightly above control on usage

Some telemetry gaps (Codex, legacy installations)

Despite this, we have high confidence in directional accuracy

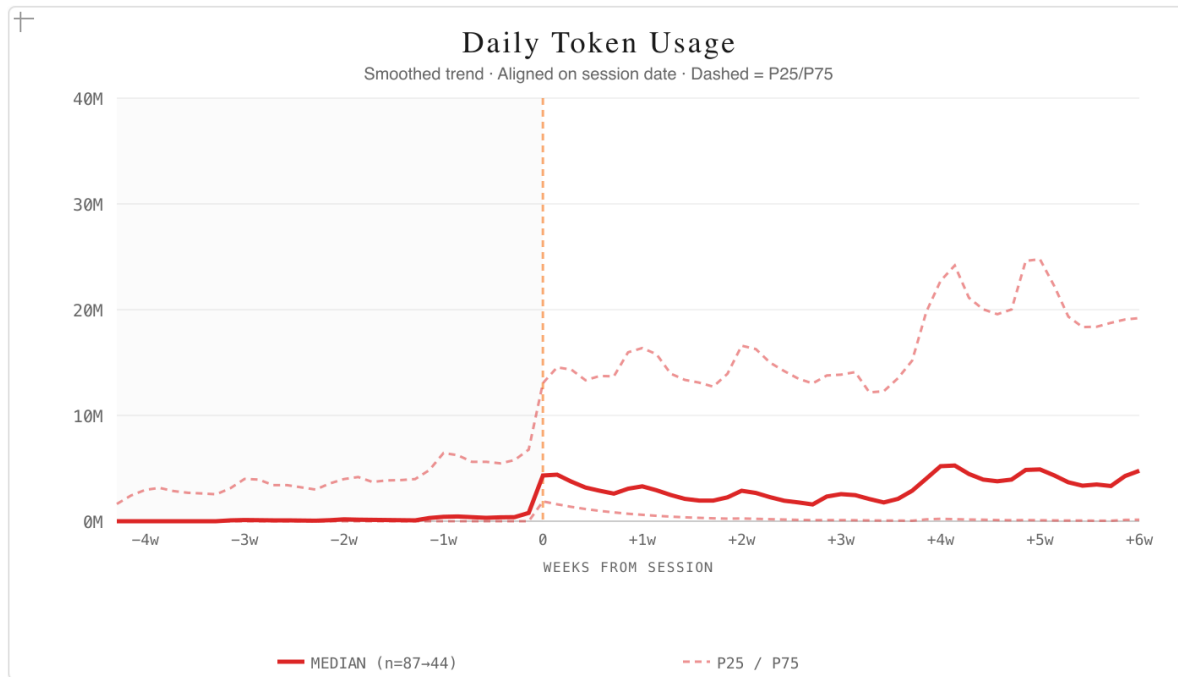
**SHOW ME
THE
RESULTS**

THE IMPACT OF THE DAY

Aligned to actual
session date, we see a
step change

No decay over six weeks

The workshop
eliminates non-users

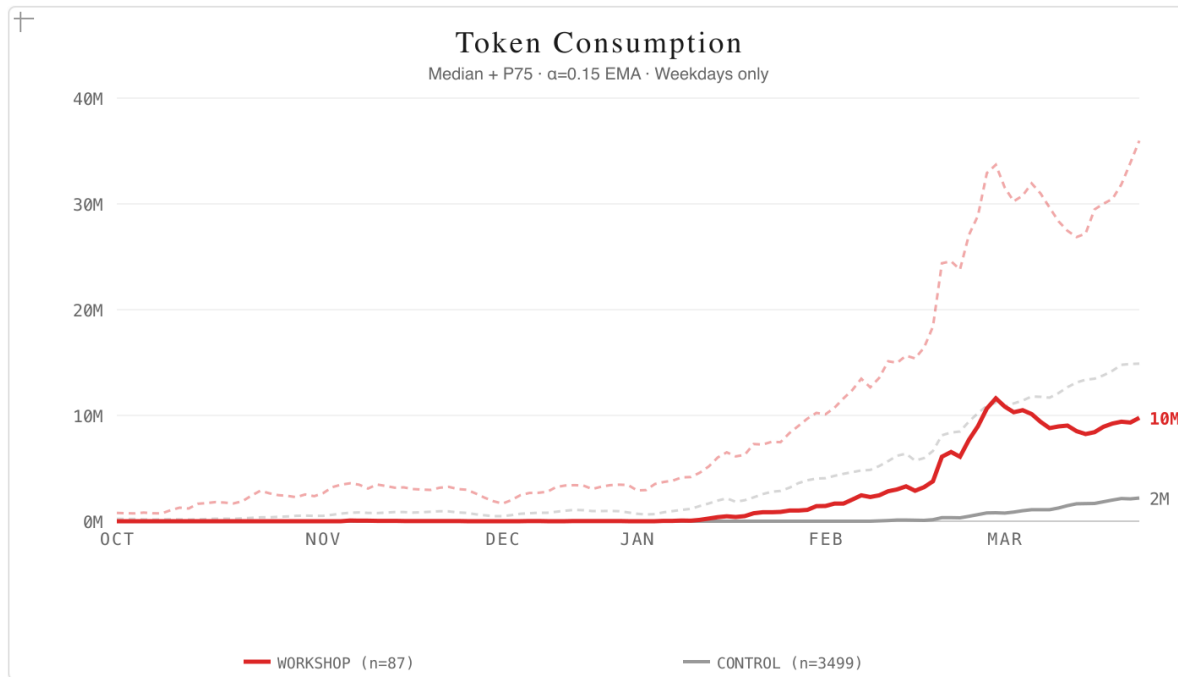


SAME PATTERN: RAW TOKENS

Both groups trending up

Workshop slope is
steeper

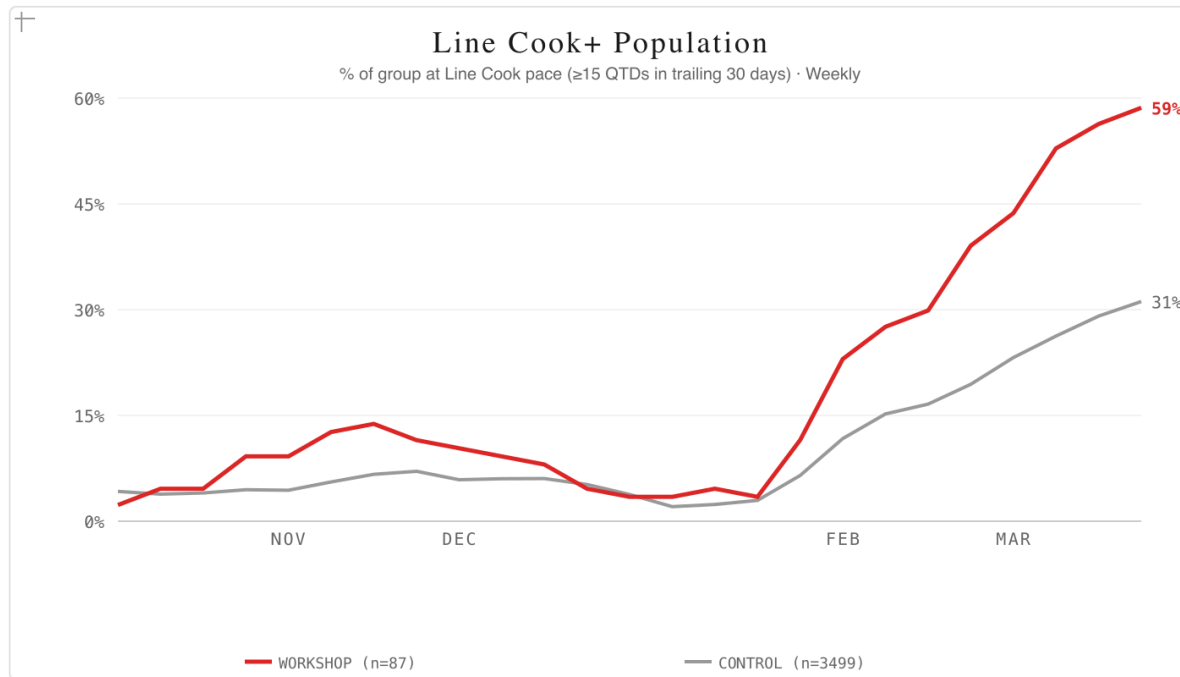
Workshop group **pulls
away** as first cohorts
complete



SAME PATTERN: LINE COOKS

**59% of workshop at
Line Cook vs 31% of
controls**

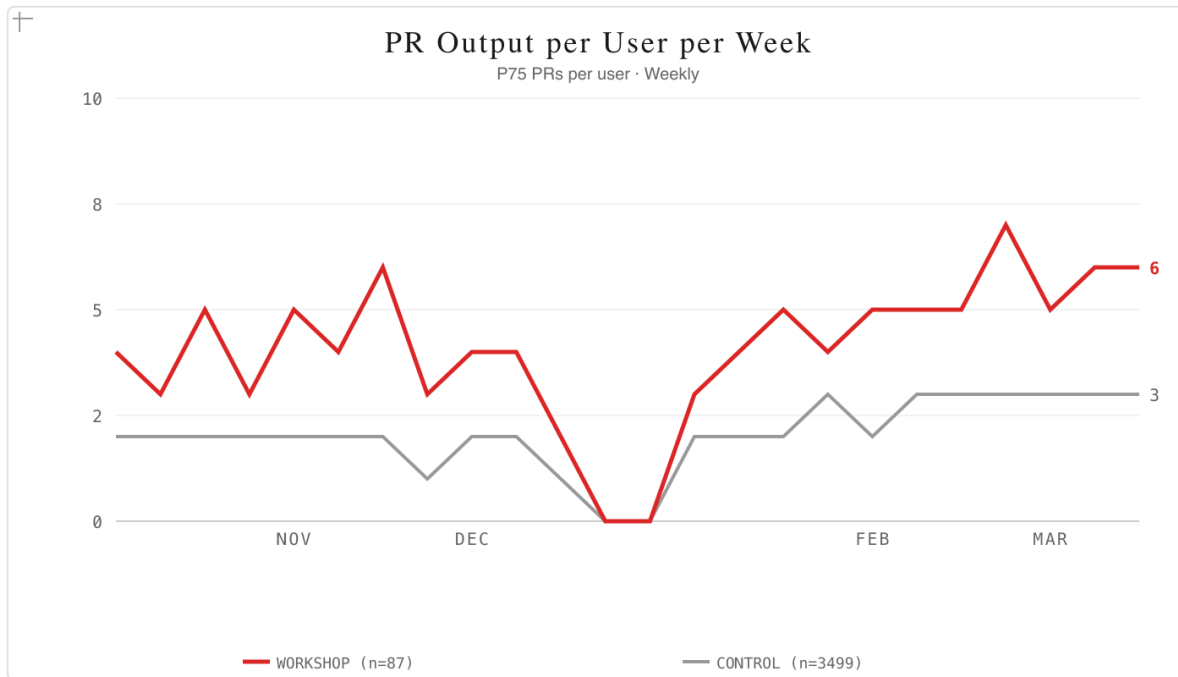
**+28pp delta, still
growing**



CODE OUTPUT: OPEN QUESTIONS

**PR count shows no
change for either group**

**Review cycle is the
bottleneck?**



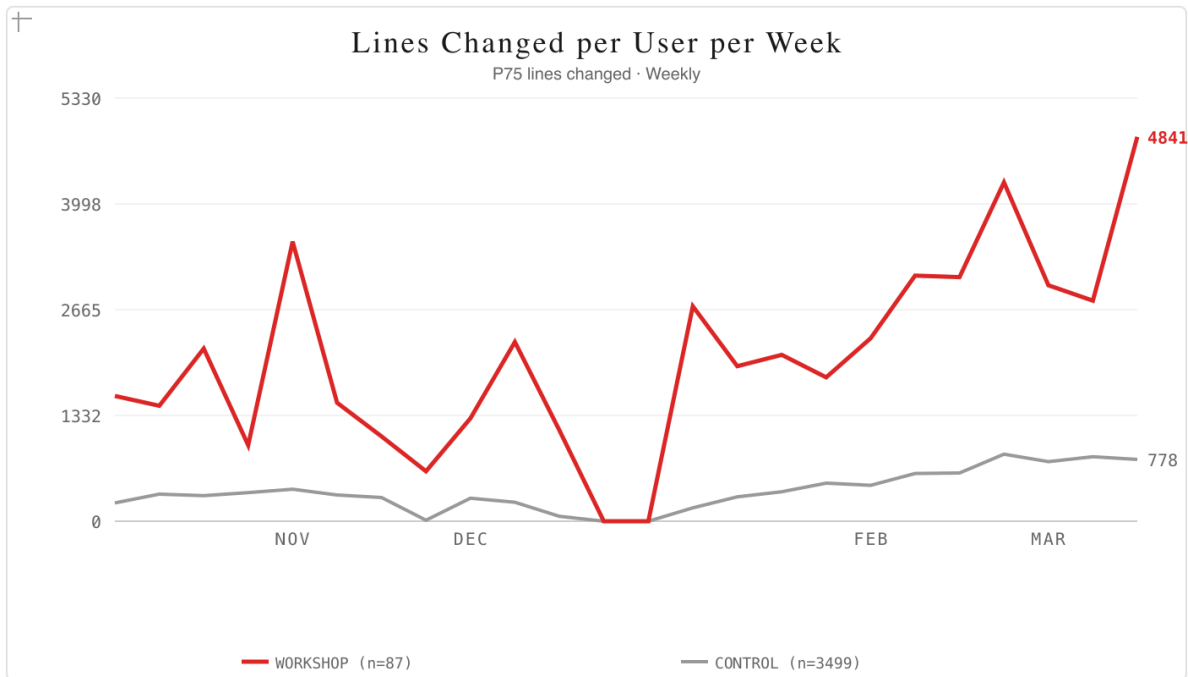
WHAT'S INSIDE THE COMMITS

**Analysis of 390 commits
by 9 authors**

Delta is entirely **more
test code**

**Longitudinal:
14% → 21%**

**Cross-sectional:
26% → 40%**

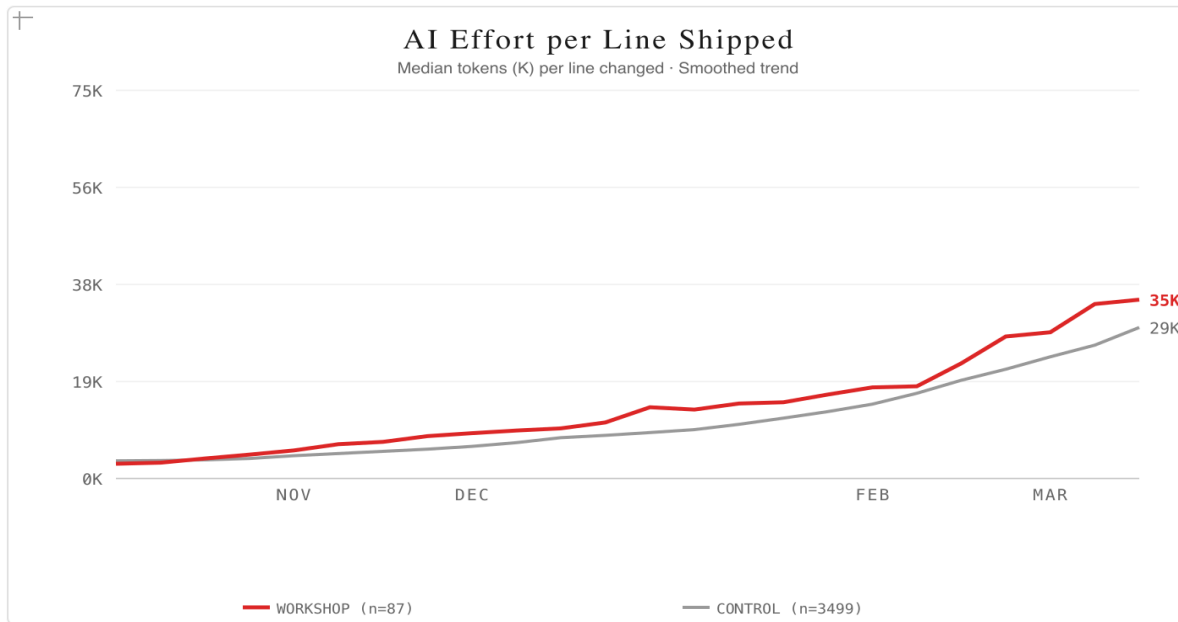


AI EFFORT PER LINE SHIPPED?

Tokens consumed per line changed — rising for all groups

More brainstorming, review, and iteration per change shipped?
Hungrier coding agents?

We aren't sure what to make of this.



HOW IT FEELS

96% would recommend the workshop

Confidence increased **+37%**

We're now running strategy at the speed we used to run bug fixes

Continuous stress markers — is this sustainable?

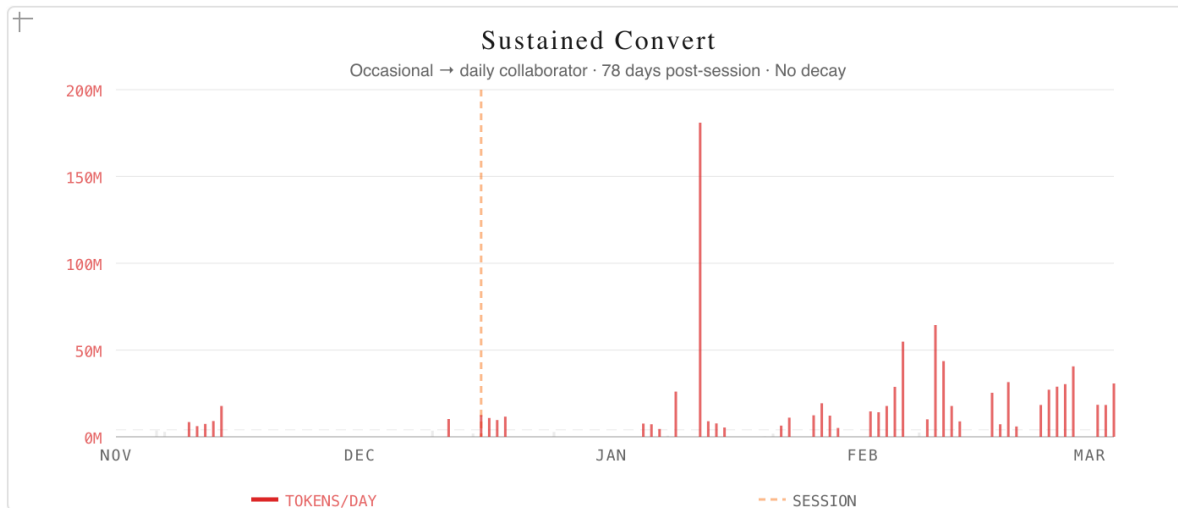


INDIVIDUAL IMPACTS

Early-career engineer -
her PR output
quadrupled.

The effect **has not**
decayed.

*"I have a PR for a bug fix I filed
months ago done within **1**
hour."*

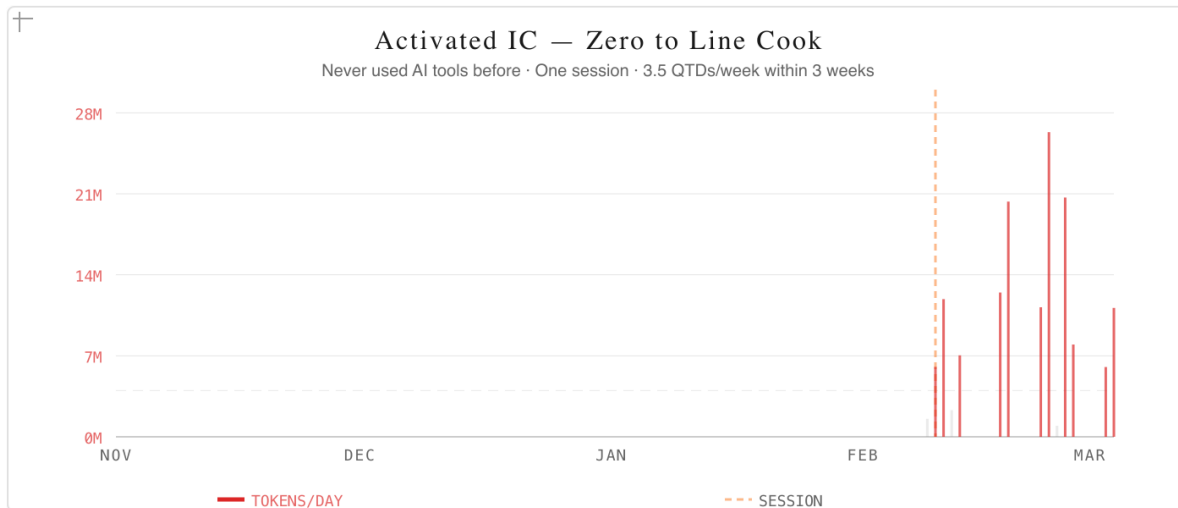


INDIVIDUAL IMPACTS

Senior engineer with 14 years of tenure

Never used an agent

Workshop → **Line Cook**

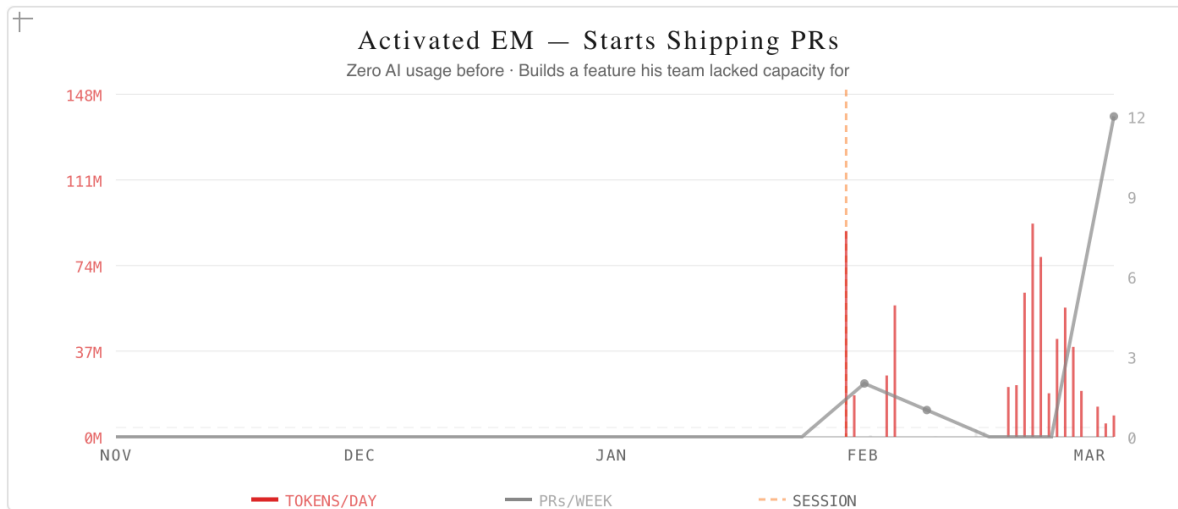


INDIVIDUAL IMPACTS

EM used to use Gemini on the web

Shipped one feature for his team, now ships tooling constantly

"Claude can take raw user requests and create a spec after brainstorming with me — **this is amazing.**"

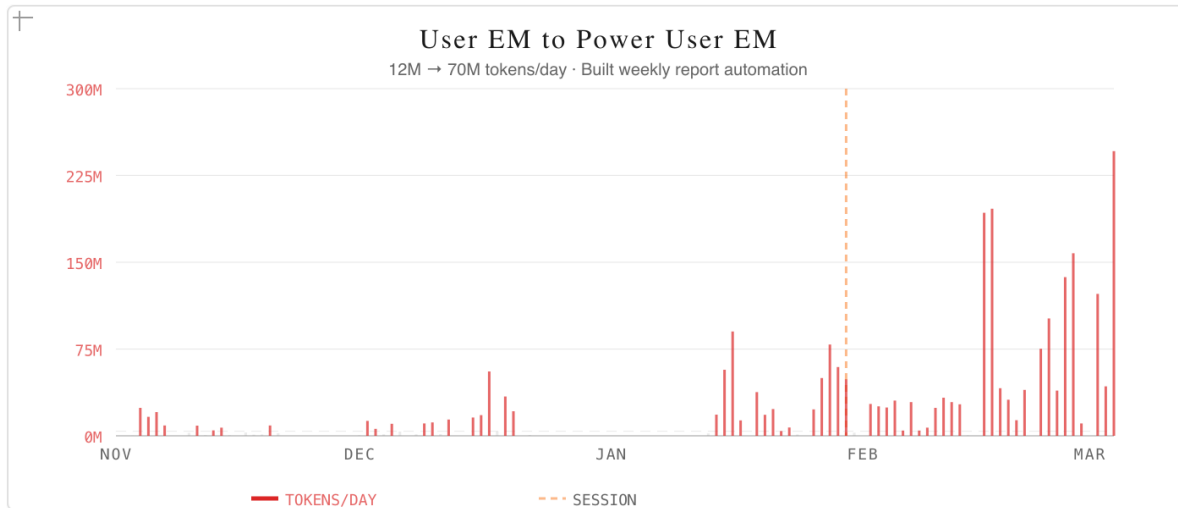


INDIVIDUAL IMPACTS

Already a power user -
went from **12M to 70M
tokens/day**

Built a status report
automation that had
taken hours manually

*"Working directly with Claude
to remove ambiguity and make
improvements was extremely
helpful."*



GENAI WEEK: **A NATURAL** **EXPERIMENT**



GENAI WEEK ATTENDEES

538 ENGINEERS

COMPANY-WIDE AI EDUCATION

EVENT
SAME HANDS-ON CONTENT

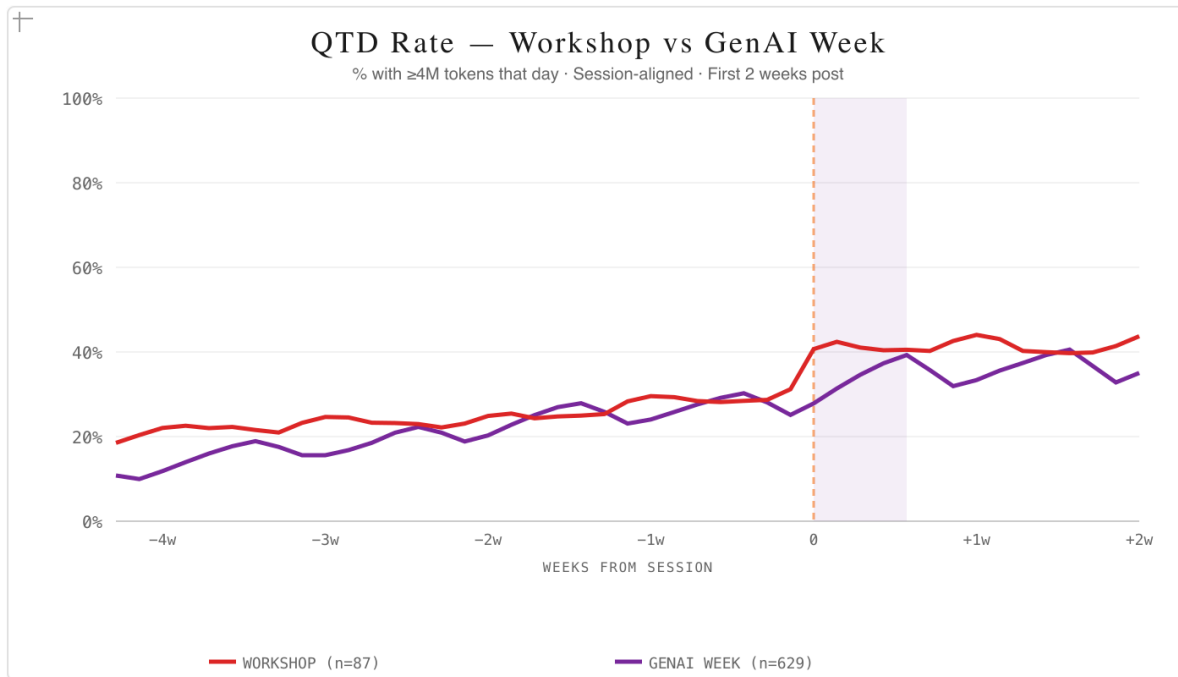
SAME TEACHERS

LARGE FORMAT CLASSES

GENAI WEEK: IMMEDIATE IMPACT

Session-aligned usage shows **immediate impact** of sessions

GenAI Week impact is **obvious** but **milder**

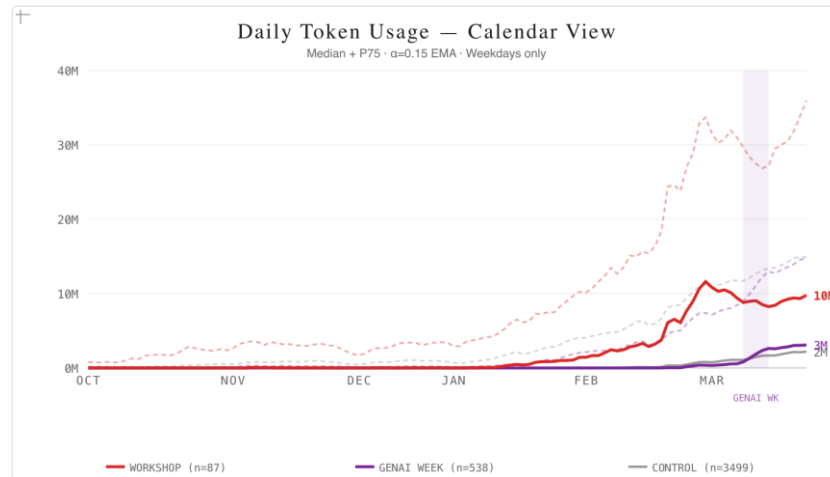
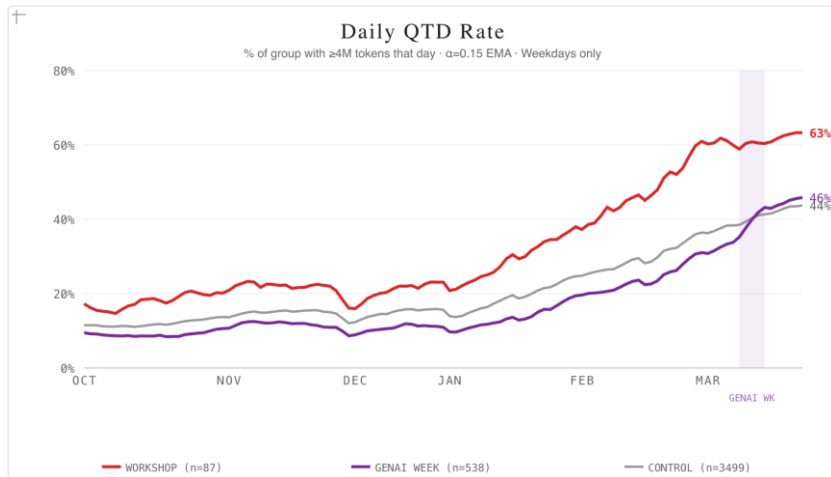


Sawtooth pattern is from GenAI Week participants sharing session dates — weekends appear at the same offset

GENAI WEEK: CATCHING UP

GenAI Week attendees start **below** the control group

During the event, they **converged**



GENAI WEEK: WHAT NEXT?

At 3 weeks post, GenAI Week group
is tracking control

Open question: will they **diverge**?

Need 4-6 more weeks of data to
answer

WHAT WE'RE WATCHING

Will the workshop group continue to diverge?

Will GenAI Week diverge from control, or plateau?

What is a reasonable expectation for QTD saturation?

Is the pace sustainable?

FOR THIS ROOM

What's the right ratio of intensive vs broad-format education?

How do you measure quality of AI-assisted work?

Who else is running controlled education programs?

THANK YOU
